

Data Driven Design of Blended Dispatching Rules

Using Preference Learning and Evolutionary Search
Case study for JSP and PFSP



Helga Ingimundardóttir
Thomas Philip Rúnarsson

University of Iceland

September 13th, 2013

Outline

Introduction

Job Shop Scheduling

Preference models

Evolutionary search with CMA-ES

Experiments

Summary and conclusions



Introduction

Job Shop Scheduling

Preference models

Evolutionary search with CMA-ES

Experiments

Summary and conclusions

Motivation

General Goal

- General goal is how to search for *good* solutions for an arbitrary problem domain.
- Automate the design of optimization algorithms.
- Use of randomly sampled problem instances and their corresponding optimal vs. suboptimal solutions.

Case Study: JSP and PFSP

Abstract

- Framework for creating dispatching rules for JSP and PFSP.
- Supervised learning based on optimal and sub-optimal solutions.
- Training data is randomly generated problem instances and their optimal solutions. Method is purely **data-driven**.
- **Linear classification** to identify good dispatches from worse ones.
- **Robust** for higher dimensions.

Keywords: Scheduling • Composite dispatching rules • JSP • PFSP
• Generating training data • Sampling • Ranking • Scalability •
Ordinal Regression • Evolutionary Search



Introduction

Job Shop Scheduling

Preference models

Evolutionary search with CMA-ES

Experiments

Summary and conclusions

Job Shop Scheduling (1)

JSP

Simple job shop scheduling problem is where n jobs are scheduled on a set of m machines, subject to constraints:

- each job must follow a predefined machine order,
- that a machine can handle at most one job at a time.

Objective: schedule the jobs so as to minimize the maximum completion time, i.e. makespan, C_{\max} .

PFSP

Permutation flow shop scheduling is the same as JSP except the predefined machine order is homogeneous for all jobs.


Job Shop Scheduling (2)

Problem space distributions used in experimental studies

type	name	size ($n \times m$)	N_{train}	N_{test}	note
JSP	$\mathcal{P}_{jrnd}^{6 \times 5}$	6×5	500	500	random
	$\mathcal{P}_{jrndn}^{6 \times 5}$	6×5	500	500	random-narrow
	$\mathcal{P}_{jrnd}^{10 \times 10}$	10×10	–	500	random
	$\mathcal{P}_{jrndn}^{10 \times 10}$	10×10	–	500	random-narrow
PFSP	$\mathcal{P}_{frnd}^{6 \times 5}$	6×5	500	500	random
	$\mathcal{P}_{frndp}^{6 \times 5}$	6×5	500	500	random-narrow
	$\mathcal{P}_{fjc}^{6 \times 5}$	6×5	500	500	job-correlated
	$\mathcal{P}_{frnd}^{10 \times 10}$	10×10	–	500	random
	$\mathcal{P}_{frndn}^{10 \times 10}$	10×10	–	500	random-narrow
	$\mathcal{P}_{fjc}^{10 \times 10}$	10×10	–	500	job-correlated

Job Shop Scheduling (3)

Simple Priority Dispatching Rules



Job Shop Scheduling (4)

Dispatching rules (DR) for constructing JSSP

- Starts with an empty schedule and adds on one job at a time.
- When a machine is free the DR inspects the waiting/available jobs and selects the job with the **highest priority**.
- Complete schedule consists of $\ell = n \times m$ sequential dispatches.
- At each dispatch k features $\phi(k)$ for the temporal schedule are calculated.
- Performance of DR is compared with its optimal makespan, as percentage relative deviation from optimality: $\rho = \frac{C_{\max}^{DR} - C_{\max}^{opt}}{C_{\max}^{opt}} \cdot 100\%$


Job Shop Scheduling (5)

Features for JSSP

ϕ	Feature description
ϕ_1	processing time for job on machine
ϕ_2	start-time
ϕ_3	end-time
ϕ_4	when machine is next free
ϕ_5	current makespan
ϕ_6	work remaining
ϕ_7	most work remaining
ϕ_8	slack time for this particular machine
ϕ_9	slack time for all machines
ϕ_{10}	slack time weighted w.r.t. number of operations already assigned
ϕ_{11}	time job had to wait
ϕ_{12}	size of slot created by assignment
ϕ_{13}	total processing time for job

Job Shop Scheduling

Example



A schedule being built at step $k = 16$. The dashed boxes represent five different possible jobs that could be scheduled next using a DR.



Introduction

Job Shop Scheduling

Preference models

Evolutionary search with CMA-ES

Experiments

Summary and conclusions

Ordinal Regression (1)

Preference learning problem

Specified by a set of **preference pairs**:

$$S = \left\{ \{z_o, +1\}_{k=1}^{\ell}, \{z_s, -1\}_{k=1}^{\ell} \mid \forall o \in \mathcal{O}^{(k)}, s \in \mathcal{S}^{(k)} \right\} \subset \Phi \times Y$$

where the set of point/rank pairs are:

- Optimal decision: $z_o = \phi^{(o)} - \phi^{(s)}$, ranked +1
- Sub-optimal decision: $z_s = \phi^{(s)} - \phi^{(o)}$, ranked -1

Ordinal Regression (2)

- Mapping of points to ranks: $\{h(\cdot) : \Phi \mapsto Y\}$ where

$$\phi_o \succ \phi_s \Leftrightarrow h(\phi_o) > h(\phi_s)$$

- The preference is defined by a linear function, i.e. **PREF model**:

$$h(\phi) = \sum_{i=1}^d w_i \phi_i = \langle w \cdot \phi \rangle.$$

- Logistic regression learns the optimal parameters w by solving:

$$\min_w \quad \frac{1}{2} \langle w \cdot w \rangle + C \sum_{j=1}^{|S|} \log \left(1 + e^{-y_j \langle w \cdot z_j \rangle} \right)$$

Generating preference set S (1)

A separate DR for each dispatch iteration

- At each dispatch k a number of data pairs are created
 - for each of the N_{train} problem instance created.
- Deliberately create a separate data set for each dispatch
 - Resulting in ℓ linear scheduling rules for solving a $n \times m$ JSSP.

Defining the size of the training set as $I = |\Phi|$, gives the size of the preference set as $|S| = 2I$.

- If I is too large, than sampling needs to be done.

Generating preference set S (2)

Previous sampling approach

The strategy was to follow some **single optimal job** $j \in \mathcal{O}^{(k)}$, thus creating $|\mathcal{O}^{(k)}| \cdot |\mathcal{S}^{(k)}|$ feature pairs at each dispatch k , resulting in a training size of:

$$I = \sum_{q=1}^{N_{\text{train}}} \left(\sum_{k=1}^{\ell} |\mathcal{O}^{(k)}| \cdot |\mathcal{S}^{(k)}| \right)$$

For the data distribution considered there, this simple sampling was sufficient for a favourable outcome. However for a considerably harder data distribution this strategy did not work well.

Generating preference set S (3)

Trajectory sampling strategies explored for S ,

S^{opt} follow some (random) optimal task

S^{cma} follow the task corresponding to highest priority, computed with fixed weights \mathbf{w} , which were obtained by optimising with CMA-ES.

S^{mwr} follow the SDR most work remaining (MWR).

S^{lwr} similar to S^{mwr} except for least work remaining (LWR).

S^{all} union of all of the above.

Generating preference set S (4)

Ranking strategies explored for S ,

S_b all optimum rankings r_1 versus all possible sub-optimum rankings $r_i, i \in \{2, \dots, n'\}$, preference pairs added.

S_f full subsequent rankings, i.e. all possible combinations of r_i and r_{i+1} for $i \in \{2, \dots, n'\}$, preference pairs added.

S_p partial subsequent rankings, i.e. sufficient set of combinations of r_i and r_{i+1} for $i \in \{2, \dots, n'\}$, preference pairs added. Note that $S_p \subset S_f$.

S_{all} union of all of the above.

where $r_1 > r_2 > \dots > r_{n'} (n' < n)$ are the rankings of the ready-list at each time step.



Introduction

Job Shop Scheduling

Preference models

Evolutionary search with CMA-ES

Experiments

Summary and conclusions

Evolutionary search

Instead of using logistic regression for to find the weights \mathbf{w} for linear preference function:

$$h(\phi) = \sum_{i=1}^d w_i \phi = \langle \mathbf{w} \cdot \phi \rangle.$$

a widely-used evolutionary algorithm, Covariance Matrix Adaptation Evolution Strategy (**CMA-ES**), is applied to directly minimise the expected relative error, i.e. $\mathbb{E}[\rho]$ (note, could also minimise $\mathbb{E}[C_{\max}]$)

Benefit No need to collect preference set S

Drawback Computationally expensive to evaluate $\mathbb{E}[\rho]$



Introduction

Job Shop Scheduling

Preference models


Evolutionary search with CMA-ES

Experiments


Summary and conclusions

Experiments (1)


Size of preference set S for $\mathcal{P}_{jrnd}^{6 \times 5}$



S_{opt}



S_{cma}




S_{mrw}

S_b in blue, S_f in red and S_p in green.

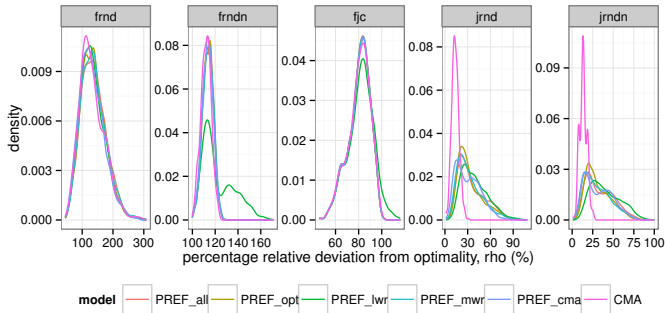
Experiments (2)

Size of preference set S_p



Experiments (3)

Linear PREF models and CMA-ES obtained weights





Introduction

Job Shop Scheduling

Preference models

Evolutionary search with CMA-ES

Experiments

Summary and conclusions

Summary and conclusions (1)

- Introduced a framework for learning linear composite dispatch rules for scheduling.
- The approaches find linear weights by either **direct optimisation with CMA-ES** or via **preference learning** by collecting preference pairs whilst sampling the state space of the schedule strategically.

Summary and conclusions (2)

CMA-ES optimisation

Benefits:

- Does not rely on optimal solutions
- Scalable

Drawbacks:

- Computationally expensive .
- Limited to linear preference function $h(\cdot)$

Future Work:

- Mediate evolutionary search by use of surrogate models which indirectly estimate mean expected error w.r.t. current population without a loss in performance

Summary and conclusions (3)

PREF models

Benefits:

- Scalable
- Robust to different data distributions

Drawbacks:

- Must know the optimal solution of the problem a priori to correctly classify optimal decisions from suboptimal ones

Future work:

- Adaptable to non-linear preferences function, i.e. project the feature space onto a higher dimension thereby updating $h(\cdot)$ to a kernel based function which should yield lower expected C_{\max}



Thank you for your attention

Questions?

Helga Ingimundardóttir, hei2@hi.is

References (1)

- [1] M. Bader-El-Den, R. Poli, and S. Fatima. Evolving timetabling heuristics using a grammar-based genetic programming hyper-heuristic framework. *Memetic Computing*, 1(3):205–219, 2009.
- [2] E. Burke, G. Kendall, J. Newall, E. Hart, P. Ross, and S. Schulenburg. Hyper-heuristics: An emerging direction in modern search technology. *International series in operations research and management science*, pages 457–474, 2003.
- [3] E. Burke, S. Petrovic, and R. Qu. Case-based heuristic selection for timetabling problems. *Journal of Scheduling*, 9:115–132, 2006.
- [4] E. K. Burke, M. Hyde, G. Kendall, G. Ochoa, E. Ozcan, and R. Qu. Hyper-heuristics: A survey of the state of the art. *Journal of the Operational Research Society (to appear)*, 2010.
- [5] E. K. Burke, M. R. Hyde, G. Kendall, and J. Woodward. Automating the packing heuristic design process with genetic programming. *Evolutionary computation*, 20(1):63–89, 2012.

References (2)

- [6] D. Corne and A. Reynolds. Optimisation and generalisation: Footprints in instance space. In R. Schaefer, C. Cotta, J. Kolodziej, and G. Rudolph, editors, *Parallel Problem Solving from Nature, PPSN XI*, volume 6238 of *Lecture Notes in Computer Science*, pages 22–31. Springer Berlin, Heidelberg, 2010.
- [7] A. S. Fukunaga. Automated discovery of local search heuristics for satisfiability testing. *Evolutionary Computation*, 16(1):31–61, 2008.
- [8] H. Ingimundardottir and T. P. Runarsson. Determining the Characteristic of Difficult Job Shop Scheduling Instances for a Heuristic Solution Method. In M. Schoenauer, editor, *Learning and Intelligent Optimization, 6th International Conference, LION 6*, Paris, 2012. Springer Lecture Notes in Computer Science.
- [9] S. Kalyanakrishnan and P. Stone. Characterizing reinforcement learning methods through parameterized learning problems. *Machine Learning*, 84(1-2):205–247, June 2011.

References (3)

- [10] X. Li and S. Olafsson. Discovering dispatching rules using data mining. *Journal of Scheduling*, 8:515–527, 2005.
- [11] E. López-Camacho, H. Terashima-Marin, G. Ochoa, and S. E. Conant-Pablos. Understanding the structure of bin packing problems through principal component analysis. *International Journal of Production Economics*, (in press):–, 2013.
- [12] A. M. Malik, T. Russell, M. Chase, and P. Beek. Learning heuristics for basic block instruction scheduling. *Journal of Heuristics*, 14(6):549–569, Dec. 2008.
- [13] S. Minton. Automatically configuring constraint satisfaction programs: A case study. *Constraints*, 1(1-2):7–43, 1996.
- [14] G. Ochoa, M. Hyde, T. Curtois, J. A. Vazquez-Rodriguez, J. Walker, M. Gendreau, G. Kendall, B. McCollum, A. J. Parkes, S. Petrovic, et al. Hyflex: a benchmark framework for cross-domain heuristic search. In *Evolutionary Computation in Combinatorial Optimization*, pages 136–147. Springer, 2012.

References (4)

- [15] S. Olafsson and X. Li. Learning effective new single machine dispatching rules from optimal scheduling data. *International Journal of Production Economics*, 128(1):118–126, 2010.
- [16] B. Pfahringer, H. Bensusan, and C. Giraud-carrier. Meta-learning by landmarking various learning algorithms. In *in Proceedings of the 17th International Conference on Machine Learning, ICML'2000*, pages 743–750. Morgan Kaufmann, 2000.
- [17] R. Poli, J. Woodward, and E. K. Burke. A histogram-matching approach to the evolution of bin-packing strategies. In *Evolutionary Computation, 2007. CEC 2007. IEEE Congress on*, pages 3500–3507. IEEE, 2007.
- [18] T. Russell, A. M. Malik, M. Chase, and P. van Beek. Learning heuristics for the superblock instruction scheduling problem. *IEEE Trans. on Knowl. and Data Eng.*, 21(10):1489–1502, Oct. 2009.

References (5)

- [19] K. Smith-Miles and L. Lopes. Generalising algorithm performance in instance space: A timetabling case study. In C. Coello, editor, *Learning and Intelligent Optimization*, volume 6683 of *Lecture Notes in Computer Science*, pages 524–538. Springer Berlin, Heidelberg, 2011.
- [20] K. Smith-Miles and T. T. Tan. Measuring algorithm footprints in instance space. *2012 IEEE Congress on Evolutionary Computation*, pages 1–8, June 2012.
- [21] J.-P. Watson, L. Barbulescu, L. D. Whitley, and A. E. Howe. Contrasting structured and random permutation flow-shop scheduling problems: Search-space topology and algorithm performance. *INFORMS Journal on Computing*, 14:98–123, 2002.
- [22] D. H. Wolpert and W. G. Macready. No free lunch theorems for optimization. *IEEE Transactions on Evolutionary Computation*, 1(1):67–82, 1997.