

Generating Training Data for Learning Linear Composite Dispatching Rules for Scheduling

Helga Ingimundardóttir
Thomas Philip Runarsson

University of Iceland

January, 2015



Outline

Introduction

Job Shop Scheduling

Preference Learning

Evolutionary search with CMA-ES

Experiments

Summary and conclusions



Introduction

Job Shop Scheduling

Preference Learning

Evolutionary search with CMA-ES

Experiments

Summary and conclusions

Motivation

General Goal

- General goal is how to search for *good* solutions for an arbitrary problem domain.
- Automate the design of optimization algorithms.
- Use of randomly sampled problem instances and their corresponding optimal vs. suboptimal solutions.

Case Study: JSP

Abstract

- Framework for creating dispatching rules for JSP.
- **Linear classification** to identify good dispatches from worse ones.
- Generate training data both from **optimal** and **suboptimal** solutions, by exploring various **trajectories** within the state-space.
- Sample training data using different **ranking** schemes.

Keywords: Scheduling • Composite dispatching rules • JSP • Generating Training Data • Trajectory Sampling Strategies • Ranking Schemes



Introduction

Job Shop Scheduling

Preference Learning

Evolutionary search with CMA-ES

Experiments

Summary and conclusions

Job Shop Scheduling (1)

JSP

Simple job shop scheduling problem is where n jobs are scheduled on a set of m machines, subject to constraints:

- each job must follow a predefined machine order,
- that a machine can handle at most one job at a time.

Objective: schedule the jobs so as to minimize the maximum completion time, i.e., makespan, C_{\max} .

Job Shop Scheduling (2)

Problem space distributions used in experimental studies

	name	size	N_{train}	N_{test}	note
JSP	$\mathcal{P}_{j.rnd}^{6 \times 5}$	6×5	500	500	random
	$\mathcal{P}_{j.rndn}^{6 \times 5}$	6×5	500	500	random-narrow

Job Shop Scheduling (3)

Dispatching rules (DR) for constructing JSP

- Starts with an empty schedule and adds on one job at a time.
- When a machine is free the DR inspects the waiting/available jobs and selects the job with the **highest priority**.
- Complete schedule consists of $\ell = n \cdot m$ sequential dispatches.
- At each dispatch k , **features** $\phi(k)$ for the temporal schedule are calculated.
- Performance of DR is compared with its optimal makespan, as percentage relative deviation from optimality: $\rho = \frac{C_{\max}^{DR} - C_{\max}^{opt}}{C_{\max}^{opt}} \cdot 100\%$


Job Shop Scheduling (4)

Features for JSP

ϕ	Feature description
ϕ_1	processing time for job on machine
ϕ_2	start-time
ϕ_3	end-time
ϕ_4	when machine is next free
ϕ_5	current makespan
ϕ_6	work remaining
ϕ_7	most work remaining
ϕ_8	slack time for this particular machine
ϕ_9	slack time for all machines
ϕ_{10}	slack time weighted w.r.t. number of operations already assigned
ϕ_{11}	time job had to wait
ϕ_{12}	size of slot created by assignment
ϕ_{13}	total processing time for job

Job Shop Scheduling (5)

Example




A schedule being built at step $k = 16$. The dashed boxes represent five different possible jobs that could be scheduled next using a DR.

Game-tree representation (1)

Example


First layer (i.e. root) – empty schedule at step $k = 1$



Game-tree representation (2)

Example


Second layer – all possible first dispatches at step $k = 2$



Game-tree representation (3)

Example

Third layer – given J_3 is dispatched first on M_3 at step $k = 3$





Introduction

Job Shop Scheduling

Preference Learning

Evolutionary search with CMA-ES

Experiments

Summary and conclusions

Ordinal Regression (1)

Preference learning problem

Specified by a set of **preference pairs**:

$$S = \left\{ \{z_o, +1\}_{k=1}^{\ell}, \{z_s, -1\}_{k=1}^{\ell} \mid \forall o \in \mathcal{O}^{(k)}, s \in \mathcal{S}^{(k)} \right\} \subset \Phi \times Y$$

where the set of point/rank pairs are:

- Optimal decision: $z_o = \phi^{(o)} - \phi^{(s)}$, ranked +1
- Suboptimal decision: $z_s = \phi^{(s)} - \phi^{(o)}$, ranked -1

and $\phi_o, \phi_s \in \Phi \subset \mathcal{F}$ are features from the collected training set Φ .

Ordinal Regression (2)

- Mapping of points to ranks: $\{h(\cdot) : \Phi \mapsto Y\}$ where

$$\phi_o \succ \phi_s \Leftrightarrow h(\phi_o) > h(\phi_s)$$

- The preference is defined by a linear function, i.e. **PREF model**:

$$h(\phi) = \sum_{i=1}^d w_i \phi = \langle w \cdot \phi \rangle.$$

- Logistic regression learns the optimal parameters w by solving:

$$\min_w \quad \frac{1}{2} \langle w \cdot w \rangle + C \sum_{j=1}^{|S|} \log \left(1 + e^{-y_j \langle w \cdot z_j \rangle} \right)$$

Generating preference set S (1)

A separate DR for each dispatch iteration

- At each dispatch k a number of data pairs are created
 - for each of the N_{train} problem instance created.
- Deliberately create a separate data set for each dispatch
 - Resulting in ℓ linear scheduling rules for solving a $n \times m$ JSP.

Defining the size of the training set as $I' = |\Phi|$, gives the size of the preference set as $|S| = I = 2I'$.

- If I is too large, than sampling needs to be done.

Generating preference set S (2)

Previous sampling approach

The strategy was to follow some **single optimal job** $j \in \mathcal{O}^{(k)}$, thus creating $|\mathcal{O}^{(k)}| \cdot |\mathcal{S}^{(k)}|$ feature pairs at each dispatch k , resulting in a training size of:

$$I' = \sum_{q=1}^{N_{\text{train}}} \left(\sum_{k=1}^{\ell} |\mathcal{O}^{(k)}| \cdot |\mathcal{S}^{(k)}| \right)$$

For the data distribution considered there, this simple sampling was sufficient for a favourable outcome. However for a considerably harder data distribution this strategy did not work well.

Generating preference set S (3)

Trajectory sampling strategies explored for adding features to Φ

- Φ^{opt} follow some (random) optimal task
- Φ^{cma} follow the task corresponding to highest priority, computed with fixed weights \mathbf{w} , which were obtained by optimising with **CMA-ES**.
- Φ^{mwr} follow the SDR most work remaining (MWR).
- Φ^{rnd} follow some random task.
- Φ^{all} union of all of the above, i.e.,

$$\Phi^{all} = \Phi^{opt} \cup \Phi^{cma} \cup \Phi^{mwr} \cup \Phi^{rnd}$$

Generating preference set S (4)

Ranking schemes implemented for adding preference pairs to S

S_b all opt rankings r_1 vs. all possible subopt rankings r_i , $i \in \{2, \dots, n'\}$

S_f full subsequent rankings, i.e., all combinations of r_i and r_{i+1} for all $i \in \{1, \dots, n'\}$.

S_p partial subsequent rankings, similar to S_f except if there are more than one operation with the same ranking, only one is needed to be compared to subsequent rank, i.e., $S_p \subset S_f$.

S_a union of all of the above, i.e.,

$$S_a = S_b \cup S_f \cup S_p$$

where $r_1 > r_2 > \dots > r_{n'}$ ($n' \leq n$) are the rankings of $\mathcal{R}^{(k)}$.



Introduction

Job Shop Scheduling

Preference Learning

Evolutionary search with CMA-ES

Experiments

Summary and conclusions

Evolutionary search

Instead of using logistic regression for to find the weights \mathbf{w} for linear preference function:

$$h(\phi) = \sum_{i=1}^d w_i \phi = \langle \mathbf{w} \cdot \phi \rangle.$$

a widely-used evolutionary algorithm, Covariance Matrix Adaptation Evolution Strategy (**CMA-ES**), is applied directly on the objective function.

Benefit No need to collect training data beforehand.

Drawback Computationally expensive to evaluate $\mathbb{E}[C_{\max}]$



Introduction

Job Shop Scheduling

Preference Learning


Evolutionary search with CMA-ES

Experiments

Summary and conclusions

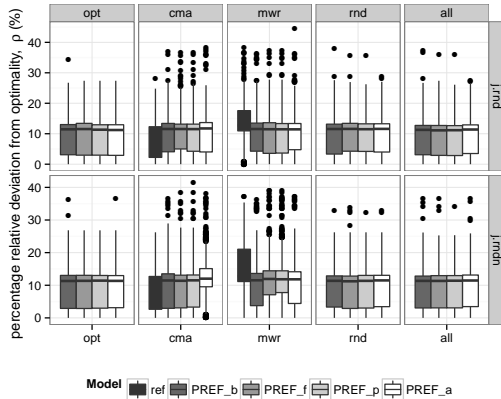
Experiments (1)

Size of preference set, $l = |S|$



Experiments (2)

Box-plot for PREF models using test set



Experiments (3)

Trajectory sampling strategies

- Learning preferences from good scheduling rules can give favourable results, e.g. Φ^{cma} and Φ^{mwr} .
- Tracking only optimal paths (Φ^{opt}) yield a generally lower mean relative error – although there was no statistical difference with Φ^{rnd} , implying the model has diverged from the learned optimal features set and is inept to determine good dispatches from that point onward.
- For $\mathcal{P}_{j.rnd}^{6 \times 5}$ the best model was based on Φ^{all} , where the suboptimal trajectories aid Φ^{opt} by adding a greater variety of preference pairs for getting out of local minima.

Experiments (4)

Results for ranking schemes

- No statistical difference between ranking schemes. However, opting for a smaller preference set then S_p is preferred.



Introduction

Job Shop Scheduling

Preference Learning

Evolutionary search with CMA-ES

Experiments

Summary and conclusions



Summary and conclusions (1)

- Introduced a framework for learning linear composite dispatch rules for scheduling.
- The approach is based on preference learning and its success is highly dependent on the preference pairs introduced to the system, i.e., the trajectories explored through the state-space. Here, optimal and suboptimal trajectories are of value.
- By partial subsequent ranking scheme it's possible to reduce the preference set without loss of performance.

Summary and conclusions (2)

Future work

- Instead of creating preference set based on collecting trajectories based on arbitrary dispatching rules (or models), it's worthwhile to continue with $PREF^{opt}$ model and collect preference set following its learned policy and use that to create a new model similar to Φ^{all} .
- Namely use the model to update itself – also known as **imitation learning**.
- Preliminary experiments have shown that a $PREF^{opt}$ with 14.07% mean relative error can be improved to 8.52% with only one unsupervised iteration, or 9.98% if iteration is supervised.¹

¹Supervised set-up: 50% chance $PREF^{opt}$ used, 50% optimal move chosen.

Thank you for your attention



HÁSKÓLI ÍSLANDS
VERKFRÆÐI- OG NATTÚRUVÍSINDASVIÐ

Questions?

Contact: Helga Ingimundardóttir, hei2@hi.is